

Chapter 1: What is an image?

Table of contents

What is an image?	1
An image is a matrix	2
Binary image	2
Grayscale image	3
Indexing and manipulating images	5
Reading and writing images	6
A note on image formats	7
Exercises	7

What is an image?

An image is a static snapshot of a scene. Let us start with an analog image and work through understanding a digital image. An example of an analog image is a photograph taken with an old fashioned camera on a physical photosensitive film. If we were to take a large magnifying glass and zoom in on a really small area, the colors on the image will change in a very ‘smooth’ manner without any discontinuities. Such an image is an analog image, which is in essence a ‘perfect’ representation of the captured object where every possible detail of an image is captured.

Our computers are extremely good at storing and manipulating massive amounts of numerical values. That is they store digital discrete values, and hence impossible for them to store numerical values for infinitesimally small intervals when moving along an image. Our computers store a ‘digitized’ version of an image. An image is digitized by dividing an analog image into grids. Each grid represents a pixel, and each pixel stores the average value of the light falling in that square, that is each pixel stores one value. The smaller the grid the smaller the area of the image capture by the grid, and hence more grids are required to cover the entire image. Making the grids smaller increased the resolution of the image, however it comes at the cost of storing more values per image (and so increasing the image size).

In a digital image, at each pixel the intensity

Obviously, all digital cameras and microscopes do not capture an analog image and convert them to a digital image, but rather they capture a digital image directly. The most common

methods used for capturing digital images in a microscope are CMOS-CCD such as used in most wide-field microscopes and PMT such as used in most confocal microscopes. The details of how these sensors work are irrelevant here, and we obtain a final digital image from them.

Now that we have divided an image into a fixed number of grids and have the average intensity at each grid, how are these values actually stored in a computer and how can we manipulate them? We will use make extensive use of two python packages,`numpy` and `scikit-image`, for our image analysis.

An image is a matrix

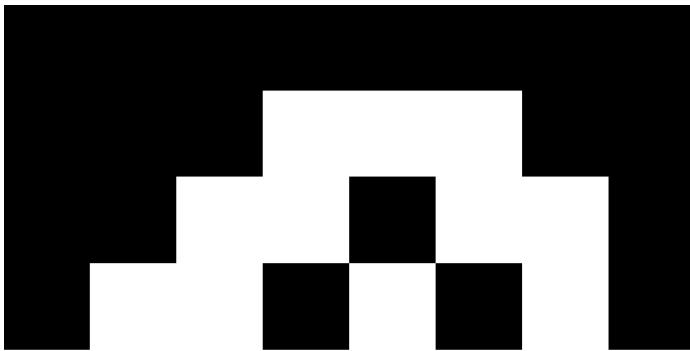
Binary image

As we have seen an image is a grid of pixels. A grid is most easily represented as matrix. The dimensions of the matrix corresponding to the size of an image. Each element of the matrix correspond to a pixel. A two dimensional matrix of size $6 * 8$ would represent an image that is 6 pixels in height and 8 pixels in width, and the intensity is stored at each pixel. The images in `ski-image` are represented as `numpy` arrays:

```
# Load the required packages
import skimage as ski
import numpy as np
import matplotlib.pyplot as plt

# A 2D array
img = np.array([[0, 0, 0, 0, 0, 0, 0, 0],
                [0, 0, 0, 1, 1, 1, 0, 0],
                [0, 0, 1, 1, 0, 1, 1, 0],
                [0, 1, 1, 0, 1, 0, 1, 0],
                [1, 1, 1, 1, 1, 1, 1, 1],
                [0, 0, 0, 0, 0, 0, 0, 0]])

# As we can view an image just like a plot
fig, ax = plt.subplots()
ax.imshow(img, cmap = 'gray')
plt.axis('off')
```



Here, all the pixels have a value of a 0 or a 1 and it represent just two colors. In a binary image, a 0 typically corresponds to a black pixel and a 1 to a white pixel (however, these are just arbitrary pseudo-colors for visualization).

Grayscale image

A grayscale image is similarly a matrix but each pixel can contain a larger range of values. The minimum value corresponds to a black pixel, the maximum value to a white pixel, a the values in between represent the distinct shades of gray. Therefore, the larger the range of values, the more shades of gray can be represented.

In the simplest case, let us consider a 2-bit image, where the smallest value at a pixel is 0 and the largest value is $2^2 - 1 = 3$. Here, 0 represents black and 3 represents white, and 1 and 2 represent shades of gray in between.

```
# A 2D image with 4 distinct values
img = np.array([[0, 0, 0, 0, 0, 0, 0, 0],
                [1, 1, 1, 1, 1, 1, 1, 1],
                [2, 2, 2, 3, 3, 2, 2, 2],
                [0, 1, 2, 3, 3, 2, 1, 0],
                [0, 1, 2, 3, 3, 2, 1, 0],
                [0, 1, 2, 3, 3, 2, 1, 0]])

# plot the image
fig, ax = plt.subplots()
ax.imshow(img, cmap = 'gray')
plt.axis('off')
```



In an 8-bit image we can have $2^8 = 256$ distinct values, with 0 being black, 255 being white, and 254 values of gray in between. Similarly a 16-bit image can have $2^{16} = 65536$ distinct values. Therefore, as the bit-depth of an image increases we can it can represent an image with more shades of gray. 16 bit images are the most typical in microscopy, though all the concepts are identical irrespective of the image type.

The figure below shows 2, 4, 8, and 16 bit images, with 0 at the top right and the maximum value at the bottom left, and each pixel increments by 1. Notice how the increment in intensity between adjacent pixels get smaller as the bit-depth of the image increases.

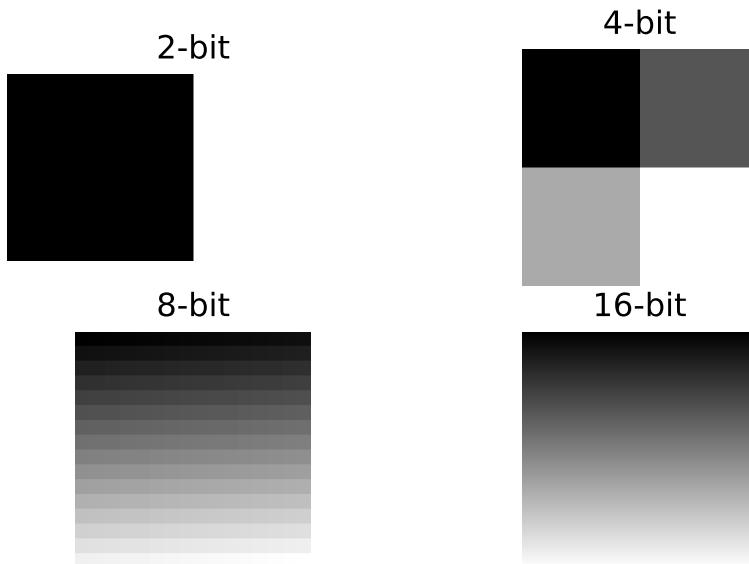
```
# A 1x2 2-bit image
img_2bit = np.arange(0, 2).reshape(1,2)
# A 2*2 4-bit image
img_4bit = np.arange(0, 4).reshape(2, 2)
# A 16*16 8-bit image
img_8bit = np.arange(0, 256).reshape(16, 16)
# A 256*256 16-bit image
img_16bit = np.arange(0, 65536).reshape(256, 256)

# plot the images
fig, ax = plt.subplots(2, 2)
ax[0, 0].imshow(img_2bit, cmap = 'gray')
ax[0, 0].set_title('2-bit')
ax[0, 0].set_axis_off()
ax[0, 1].imshow(img_4bit, cmap = 'gray')
ax[0, 1].set_title('4-bit')
ax[0, 1].set_axis_off()
ax[1, 0].imshow(img_8bit, cmap = 'gray')
ax[1, 0].set_title('8-bit')
ax[1, 0].set_axis_off()
```

```

ax[1, 1].imshow(img_16bit, cmap = 'gray')
ax[1, 1].set_title('16-bit')
ax[1, 1].set_axis_off()

```



Indexing and manipulating images

Since an image is represented as a matrix, we can similarly manipulate index and manipulate an image just like manipulating any matrix. The image coordinate [0, 0] is at the top left of the image in `numpy`. As the x-coordinate increases we go down the image, and as the y-coordinate increases we go the right of the image. A few examples to get started with image indexing are shown below.

```

# Create an 8x8 image with all 0s
img = np.zeros((8, 8))

# Index a specific pixel by specifying the both the x and y coordinates
# top left pixel
img[0][0] = 1
# bottom right pixel
img[7][7] = 1

# Index a range of pixels by specifying the start and stop coordinates
# for both x and y coordinates. The coordinates are specified in
# format `start:end-1`.
# Set the top right of the image
img[0:2, 6:8] = 2

# We can access all the columns by specifying ':' as the column index

```

```

















































































































































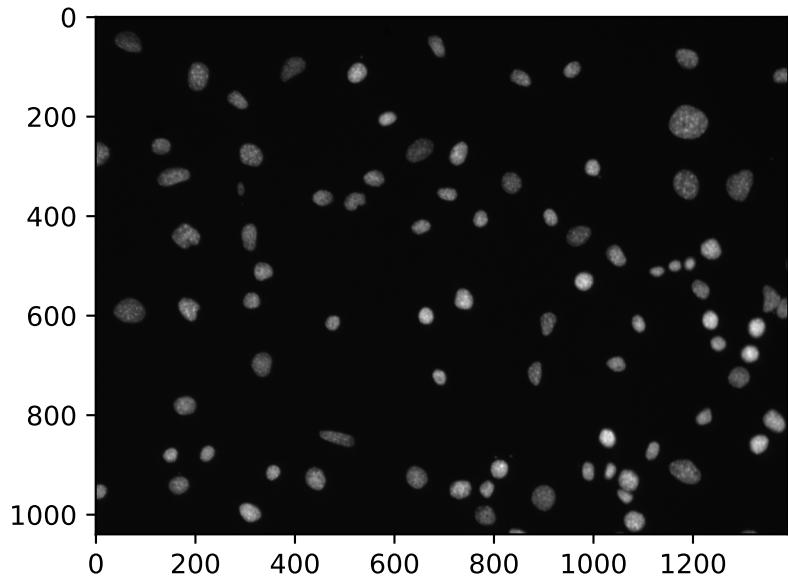



<img alt="A 5x5 grid of colored squares. The colors are arranged in a pattern: top-left square is dark gray, followed by a black row, then a light gray row, then a black row, and
```

```
# Write the manipulated image to file.  
ski.io.imsave('data/out.tif', img, check_contrast = False)
```

(1040, 1392)

uint16



A note on image formats

Exercises